

# Phase 1: Notes and Hints

- Getting started
  - USLOSS is a library: **libusloss.a**. It is located in **/home/cs452/fall115/usloss/linux**.
    - The USLOSS source is available at **/home/cs452/fall115/usloss/usloss-2.9.tgz**
  - Your phase1 code will be compiled into a library named: **libphase1.a** in your directory.
  - To execute a test case, you link the **.o** file of the test case with the **libusloss.a** and **libphase1.a** libraries.
    - Typing '**make**' will create the **libphase1.a** library in your directory.
    - Typing '**make test00**', for example, will create an executable test case named '**test00**' in your directory.
- Header files for Phase 1:
  - **/home/cs452/fall115/phase1/phase1.h**: contains function prototypes & constants to be used in this phase.
  - **/home/cs452/fall115/usloss/linux/include/usloss.h**: contains function prototypes for USLOSS library functions, many useful constants.
  - **kernel.h**: your data structures and constants for phase1.

- Mode and interrupts
  - All functions in this phase have to be executed in kernel mode.
  - Test for kernel mode must be done for each phase 1 function, since processes must be in kernel mode to call these.
    - See discussion of Process Status Register (p. 3) and functions to access this register (p. 9) in the USLOSS manual.
    - Some useful constants are defined in **usloss.h**.
  - Enabling and disabling interrupts
    - Manipulate the appropriate bits in the PSR register.
    - Interrupts can be turned on and off only in kernel mode
  - When should interrupts be disabled? (A key point in writing correct phase 1 functions.)

- **startup()** function
  - Called by USLOSS. Note: **main()** is inside the USLOSS library, not in your phase 1 code.
  - Initialize your kernel data structures, in particular, the process table.
    - The **struct procStruct** in the provided **kernel.h** is not complete. Add/change fields as needed!
  - Ready List(s): for processes waiting to run. Choose one of:
    - Single queue arranged by priorities; or
    - Multiple queues, one per priority.
- Sentinel process:
  - Call **fork1()** to create the sentinel process.
  - Sentinel runs when there are no more runnable processes
    - It is the only process with priority 6; its status is always READY.
  - The provided **sentinel()** function in **skeleton.c** is complete.
    - You will need to provide the **checkDeadlock()** function that **sentinel()** calls.
- Call **fork1()** to create a process that will execute **start1()**.
  - **start1()** will be inside every phase 1 test case.
  - **start1()** will be the beginning of phase 2.

**fork1()** function:

- Test for kernel mode.
- Initialize process table entry, which will include:
  - Check for valid stack size and priorities; there is a minimum stack size.
    - Use **malloc()** (or **posix\_memalign()**) to create the stack. Save stack information in process table so stack can be freed in **quit()**.
      - This will be the only call to **malloc** (or **posix\_memalign**) in phases 1, 2, 3, and 4!!
      - You will get to dynamically allocate memory again in phase 5.
- Store function pointer and argument value:
  - **fork1()** does not execute the function.
  - **launch()** does.
- Initialize context using the USLOSS **USLOSS\_ContextInit()** function.
  - Use **launch()** as the function pointer passed to **USLOSS\_ContextInit()**.
  - Necessary: must enable interrupts before starting function, and correctly handle return from function.
    - Unix analogy: What happens if **main** calls **return** instead of **exit**?
- Call the **dispatcher()** — let the **dispatcher()** decide whether the parent or the child runs next!
- Enable interrupts (for the parent) and return the pid of the child process.

**join()** function: There are three cases:

- The process has no children.
- Child(ren) **quit()** before the **join()** occurred.
  - Return the pid and quit status of one quit child and finish the clean up of that child's process table entry.
  - Report on quit children in the order the children **quit()**.
- No (unjoined) child has **quit()** ... must wait.
  - How?
  - After wait is over: return the pid and quit status of the child that quit.
    - Where does the parent find these?
- The child status that join returns is the argument that the child passed to **quit()**.

### **quit()** function

- Error if a process with active children calls **quit()**. Halt USLOSS with appropriate error message.
- Cleanup the process table entry (but not entirely, see **join()**)
  - Two cases:
    - Parent has already done a **join()**, or
    - Parent has not (yet) done a **join()**.
- Unblock processes that **zap**'d this process (see below).
- May have children who have **quit()**, but for whom a **join()** was not (and now will not) be done.
  - This is not an error.

### **dispatcher()** function: Decides which process gets to run.

- When is the dispatcher called?
- Checks if the current process can continue running: Has it been time-sliced? Has it been blocked? Is it still the highest priority among READY processes?
- Select a new process and perform a context switch in order to get it running.
- Choose according to scheduling policy: see Section 3.2 of Phase 1 handout.

Clock interrupt handler (see Section 3.2 of USLOSS manual).

- Defer working on this until your **fork1**, **join**, **quit**, and **dispatcher** functions are all working.
- Interrupt vector is defined by USLOSS as an array of pointers to void functions with 2 integer arguments:  

```
extern void (*USLOSS_IntVec[NUM_INTS]) (int dev, int unit); /* from usloss.h */
```
- Initialize the appropriate slot to point to your clock handler function.
- The other slots can be ignored (they will show up again in later phases).
- **clockHandler()** function.
  - Checks if the current process has exceeded its time slice. Calls **dispatcher()** if necessary.
    - Time slice is 80 ms (milliseconds).
  - The **USLOSS\_clock()** function returns time in microseconds (= 1,000 ms); thus, time slice is 80,000  $\mu$ s.

**zap()** and **isZapped()** functions.

- The zap'ing process blocks until the zap'd process quits.
- Store information in the process table of both the zap'ing and zap'd processes.
- Zap'ing process blocks itself. (How?)
- The zap'd process, when it quits, needs to know about the zap'ing process.
- Note: there can be more than one zap'ing process! During a **quit()**, unblock all zap'ing processes.

```

TARGET = libphase1.a
ASSIGNMENT = 452phase1
CC = gcc
AR = ar

COBJS = phase1.o p1.o
CSRCS = ${COBJS:.o=.c}

HDRS = kernel.h

INCLUDE = ./usloss/include

CFLAGS = -Wall -g -I${INCLUDE}
#CFLAGS += -D_XOPEN_SOURCE      # use for Mac, NOT for Linux!!

LDFLAGS = -L. -L./usloss/lib

TESTDIR = testcases
TESTS = test00 test01 test02 test03 test04 test05 test06 test07 test08 \
        test09 test10 test11 test12 test13 test14 test15 test16 test17 \
        test18 test19 test20 test21 test22 test23 test24 test25 test26
LIBS = -lphase1 -lusloss

$(TARGET): $(COBJS)
$(AR) -r $@ $(COBJS)

$(TESTS): $(TARGET) $(TESTDIR)/$$@.c p1.o
$(CC) $(CFLAGS) -I. -c $(TESTDIR)/$@.c
$(CC) $(LDFLAGS) -o $@ $@.o $(LIBS) p1.o

```



```
clean:
    rm -f $(COBJS) $(TARGET) test??.o test?? core term*.out p1.o

phase1.o:    kernel.h

submit: $(CSRCS) $(HDRS) Makefile
    tar cvzf phase1.tgz $(CSRCS) $(HDRS) Makefile
```

Selected parts of `/home/cs452/summer05/phase1/skeleton.c`

```
void startup()
{
    int i;          /* loop index */
    int result; /* value returned by call to fork1() */

    /* initialize the process table */

    /* Initialize the Ready list, etc. */
    if (DEBUG && debugflag)
        USLOSS_Console("startup(): initializing the Ready & Blocked lists\n");
    ReadyList = NULL;

    /* Initialize the clock interrupt handler */

    /* startup a sentinel process */
    if (DEBUG && debugflag)
        USLOSS_Console("startup(): calling fork1() for sentinel\n");
    result = fork1("sentinel", sentinel, NULL, USLOSS_MIN_STACK, SENTINELPRIORITY);
    if (result < 0) {
        if (DEBUG && debugflag)
            USLOSS_Console("startup(): fork1 of sentinel returned error, halting...\n");
        USLOSS_Halt(1);
    }
}
```

```

/* start the test process */
if (DEBUG && debugflag)
    USLOSS_Console("startup(): calling fork1() for start1\n");
result = fork1("start1", start1, NULL, 2 * USLOSS_MIN_STACK, 1);
if (result < 0) {
    USLOSS_Console("startup(): fork1 for start1 returned an error, halting...\n");
    USLOSS_Halt(1);
}

USLOSS_Console("startup(): Should not see this message! ");
USLOSS_Console("Returned from fork1 call that created start1\n");

return;
} /* startup */

int sentinel (char *dummy)
{
    if (DEBUG && debugflag)
        USLOSS_Console("sentinel(): called\n");
    while (1)
    {
        checkDeadlock();
        USLOSS_WaitInt();
    }
} /* sentinel */

```

```

int fork1(char *name, int (*procCode)(char *), char *arg, int stacksize, int priority)
{
    int proc_slot;

    if (DEBUG && debugflag)
        USLOSS_Console("fork1(): creating process %s\n", name);

    /* test if in kernel mode; halt if in user mode */

    /* Return if stack size is too small */

    /* find an empty slot in the process table */

    /* fill-in entry in process table */
    if ( strlen(name) >= (MAXNAME - 1) ) {
        USLOSS_Console("fork1(): Process name is too long.  Halting...\n");
        USLOSS_Halt(1);
    }
    strcpy(ProcTable[procSlot].name, name);
    ProcTable[procSlot].start_func = f;
    if ( arg == NULL )
        ProcTable[proc_slot].start_arg[0] = '\0';
    else if ( strlen(arg) >= (MAXARG - 1) ) {
        USLOSS_Console("fork1(): argument too long.  Halting...\n");
        USLOSS_Halt(1);
    }
    else
        strcpy(ProcTable[procSlot].start_arg, arg);
}

```

```

/* Initialize context for this process, but use launch function pointer for
 * the initial value of the process's program counter (PC)
 */
USLOSS_ContextInit(&(ProcTable[proc_slot].state), USLOSS_PsrGet(),
                  ProcTable[procSlot].stack, launch);

/* for future phase(s) */
pl_fork(ProcTable[proc_slot].pid);

} /* fork1 */

void launch()
{
    int result;

    if (DEBUG && debugflag)
        USLOSS_Console("launch(): started\n");

    /* Enable interrupts */
    enableInterrupts();

    /* Call the function passed to fork1, and capture its return value */
    result = Current->start_func(Current->start_arg);

    if (DEBUG && debugflag)
        USLOSS_Console("Process %d returned to launch\n", Current->pid);

    quit(result);
} /* launch */

```