

# Phase 2: Notes and Hints

- Getting started
  - Provided starter files and test cases for Phase 2 are located in: `~cs452/fall15/phase2`.
  - USLOSS is a library: `libusloss.a`. It is located in `~cs452/fall15/lib`.
  - Your phase2 code will be compiled into a library named: `libphase2.a` in your directory.
  - To execute a test case, you link the `.o` file of the test case with the `libusloss.a` and `libphase1.a` libraries.
    - Typing `'make'` will create the `libphase2.a` library in your directory.
    - Typing `'make test00'`, for example, will create an executable test case named `'test00'` in your directory.
  - You can use your own `libphase1.a` library:
    - In the provided *Makefile*, use the line:  
`PHASE1LIB = phase1`
    - a copy of your `libphase1.a` will need to be in your phase 2 directory for this to work
  - You can use Patrick's `libpatrickphase1.a` library:
    - In the provided *Makefile*, use the line:      In the provided *Makefile*, use the line:  
`PHASE1LIB = patrickphase1`                      `PHASE1LIB = patrickphase1debug`  
to link to a phase 1 library with debugging output from phase 1.
  - Your phase 2 will be graded using Patrick's `libpatrickphase1.a` library.

- Header files for Phase 2:
  - *~cs452/fall15/include/phase2.h* and *~cs452/fall15/include/phase1.h*: contain function prototypes and constants to be used in this phase.
  - *~cs452/fall15/include/usloss.h*: contains function prototypes for USLOSS library functions, many useful constants.
  - *message.h*: your data structures and constants for phase2.
- Mode and interrupts
  - All functions in phase 2 have to be executed in kernel mode.
  - Test for kernel mode must be done for each phase 2 function, since processes must be in kernel mode to call these.
  - Enabling and disabling interrupts
    - Manipulate the appropriate bits in the PSR register.
    - Interrupts can be turned on and off only in kernel mode
  - When should interrupts be disabled? (A key point in writing correct phase 2 functions.)
  - Remember that calls to **unblock\_proc** and **block\_me** will result in interrupts being enabled!

- **start1** function:
  - The phase 1 library will use **fork1** to create a process at priority 1 that will execute the **start1** code that you provide in phase 2. Thus, **start1** is the entry point for phase 2. When the code in **start1** starts executing, there will be two processes already created: **sentinel** and **start1**.
  - Initialize your phase 2 data structures, in particular, the mailbox and mail slot arrays, and the phase 2 process table.
    - The **struct mailbox** and **struct mail\_slot** structures in the provided *message.h* are not complete. Add/change fields as needed.
    - You will need a process table for Phase 2. You cannot modify/extend the phase 1 process table! Use **MAXPROC** for the size of the phase 2 process table.
      - Note: When using Patrick's *libpatrickphase1.a* code, you can use **getpid() % MAXPROC** to determine which slot in your phase 2 process table to use.
  - Create I/O mailboxes and initialize the **int\_vec** and **sys\_vec** arrays (see below).
  - Use **fork1** to create the test process: **start2**.
    - priority 1.
    - stack size = **4 \* USLOSS\_MIN\_STACK**.
  - **start1** should then block on a **join** to wait for **start2** to quit.

### Mailbox functions and notes:

- **MboxCreate:**
  - Allocate and initialize a location in your mailbox array. **MAXMBOX** from *phase2.h* is the size of this array.
  - Do not allocate slots for the mailbox at this time — they are allocated only as needed to hold messages.
- **MboxSend** basics:
  - Check for possible errors (message size too large, inactive mailbox id, etc.).
    - Return -1 if errors are found.
  - If a slot is available in this mailbox, allocate a slot from your mail slot table. **MAXSLOTS** determines the size of this array. **MAX\_MESSAGE** determines the max number of bytes that can be held in the slot.
    - Note: if the mail slot table overflows, that is an error that should halt USLOSS.
    - Further note: for conditional send, the mail slot table overflow does not halt USLOSS.
      - Return -2 in this case.
  - Copy message into this slot. Use **memcpy**, not **strcpy**: messages can hold any type of data, not just strings.
  - Block the sender if this mailbox has no available slots. For example, mailbox has 5 slots, all of which already have a message.

- **MboxReceive** basics:
  - Check for possible errors (inactive mailbox id, etc.).
    - Return -1 in this case.
  - If one (or more) messages are available in the mailbox: **memcpy** the message from the slot to the receiver's buffer.
  - Free the mailbox slot.
  - Block receiver if there are no messages in this mailbox.
- Basic tests:
  - *test00.c* — *test03.c*: Create **start2**, create mailbox(es), maximum number of mailboxes test.
  - *test04.c*: Creates two processes. Higher priority process sends to a mailbox. Lower priority process receives from mailbox.
  - *test05.c*: Creates two processes. Higher priority process receives (and will block since no message has yet been sent). Lower priority process then sends, unblocking higher priority process.
  - *test06.c*: Same as *test04.c*, but 5 messages are sent instead of 1. All should fit, since mailbox has 5 slots.

### Interrupt handlers and i/o mailboxes:

- **MboxCreate**, **MboxCondSend**, and **MboxReceive** must be working before starting work on these!
- Will need to create i/o mailboxes in **start1** for the following devices:
  - Clock: 1 mailbox.
  - Terminals: 4 mailboxes, one for each terminal unit.
  - Disks: 2 mailboxes, one for each disk unit.
- New version of **clock\_handler** needed for phase 2.
  - Error check: is the device actually the clock device?
  - Still must handle time slicing: call the phase 1 **time\_slice** function when necessary.
  - Must conditionally send to the clock i/o mailbox every 5th clock interrupt.
- Disk and terminal handlers.
  - Error checks: Is the device the correct device? Is the unit number in the correct range?
  - Read the device's status register by using the **USLOSS\_DeviceInput** function.
  - Conditionally send the contents of the status register to the appropriate i/o mailbox.
    - Conditional send is used so the low-level device handler is never blocked on the mailbox.

- **waitdevice**
  - Process calls this when it wants to receive results of i/o operations (see *test13.c* and *test14.c*).
  - Use **MboxReceive** on the appropriate mailbox, as indicated by the parameters to **waitdevice**.
- **check\_io**
  - Phase 1 needs to have the following test done in **checkDeadlock** (is already in *libpatrickphase1.a*):
 

```
if ( check_io() == 1 )
    return;
```
  - This will cause the sentinel to call the function **USLOSS\_waitInt** when one, or more, processes are blocked on i/o mailboxes.
  - You will need to provide **check\_io** as part of your phase 2 code.
    - Return 1: if at least one process is blocked on an i/o mailbox.
    - Return 0: otherwise.

- Other items: Ask us about these when you get to them (that is, after you have the other stuff above working)!
  - Releasing mailboxes that have blocked processes.
  - Priority inversion when several blocked senders (or receivers) are present at different priorities.
  - Zero-slot mailboxes.
  - **systemCallVec[ ]** and system calls.